

Intro to Regular Expressions

Richard Mills

Outline

Introduction

Working with file-names

Using patterns within code

Wrap up

Introduction

What are regular expressions?

- Also known as *regex*.
- The concept dates back to the 1950s.
- A sequence of characters that define a search *pattern* for text.
- Frequently used for:
 - Checking to see if part/all of a string meets some criteria.
 - Perform (advanced) find-and-replace operations within a string.
 - Commonly used in compilers/interpreters for parsing user code.
 - Input validation for websites.
 - Frankly... Useful any time when working with text data!

Wildcards

- Certain Excel formulae allow the use of *wildcards* (ie. regex-lite) when matching against text.
 - Count number of strings with only 3 characters:
`=COUNTIF(A1:A100, "???)`
 - Count all strings beginning with ABC:
`=COUNTIF(A1:A100, "ABC*")`
 - ? represents any single character.
 - * represents zero or more of any character.
- Other supporting functions include MATCH, VLOOKUP, SUMIF and SUMIFS.
- There is similar functionality in Access via the LIKE operator.
For example... `TABLE_NUMBER IS LIKE "73?"`

Working with file-names

What's the issue?

- Suppose we have the following files within a folder:

TOP_SECRET_DATA_01-03-2020.CSV

TOP_SECRET_DATA_20-03-2020.CSV

TOP_SECRET_DATA_25-04-2020.CSV

- The date within the file-name gives us the effective date of the data within.
- **How could we extract those dates?**

What's the issue?

- What about now?

TOP_SECRET_DATA_01-03-2020.CSV

TOP_SECRET_DATA_V2_20-03-2020.CSV

TOP_SECRET_DATA_Adj_25-04-2020.CSV

Possible solution

If we were *particularly* determined, we could use the following pseudo-code:

```
for i from 1 to LEN(string)-10:
    if is_digit(string[i]) and is_digit(string[i+1]):
        if string[i+2] = "-":
            if is_digit(string[i+3]) and is_digit(string[i+4]):
                if string[i+5] = "-":
                    if is_digit(string[i+6]) .... :
                        date_part = string[i:i+10]
                        return date_part

error 'No date!'
```

Is there an alternative way?

Possible solution

- In *regex speak* we are looking for the following pattern:

`\d\d-\d\d-\d\d\d\d`

- Where each `\d` corresponds to a single digit (0-9).
- This would match any of the dates contained within the above file names.
- However, it *could* leave to false positives, eg:
 - RANDOM_00-01-0002.CSV
 - 1234-56-78901234.CSV
- **Can we refine it?**

Possible solution

- Previously we used `\d` to match a single digit.
- We could instead use a **character class** to restrict this behaviour.
- This allows us to specify a list (or range) of permitted characters.
- As an example, we'll assume that:
 - The *day* part can start with any of 0, 1, 2 or 3.
 - The *month* part can start with either 0 or 1.
 - The *year* part will start with either 19 or 20.
- We can update our pattern to be:

`[0-3]` \d- `[01]` \d- `(19|20)` \d\d

Possible solution

`[0-3]\d-[01]\d-(19|20)\d\d`

- Note the use of the following elements:
 - `[0-3]` which matches against any of 0, 1, 2 or 3.
 - `[01]` which matches either 0 or 1.
 - `(19|20)` which matches either 19 or 20.

Note the use of `(` and `)` above.

Possible solution

`[0-3]\d-[01]\d-(19|20)\d\d`

- Note the 2x `\d` at the end.
- We could instead use a **quantifier** to remove duplication:

`[0-3]\d-[01]\d-(19|20)\d{2}`

- Both approaches are equivalent!

Possible solution

`[0-3]\d-[01]\d-(19|20)\d{2}`

- What if the day and month components didn't always have a leading 0?
 - TOP_SECRET_DATA_1-03-2020.CSV
 - TOP_SECRET_DATA_V2_20-3-2020.CSV
 - TOP_SECRET_DATA_Adj_5-4-2020.CSV
- We could use `?` which is another type of quantifier.
 - This matches the prior element 0 **or** 1 times.

`[0-3]?\d-[01]?\d-(19|20)\d{2}`

Possible solution

- The last change we *might* want to make is to ensure that either end of a potential date does **not** touch a digit.
- For this, we can use:
 - `\D` to match a non-digit character.
 - `(?<=...)` to match an element immediately preceding our date.
 - `(?=...)` to match an element immediately after our date.
- Our final proposed pattern is:

`(?<=\D)` `[0-3]? \d- [01]? \d- (19|20) \d{2}` `(?=\D)`

Possible solution

①	②	③	④	⑤	⑥	⑦	⑧	⑨
(?<=\D)	[0-3]?	\d	-	[01]? \d	-	(19 20)	\d{2}	(?=\D)

- 1 The character immediately before our date must be a non-digit
- 2 The **day** part must start with an optional 0, 1, 2 or 3
- 3 ... followed by a 0–9
- 4 ... followed by a –
- 5 ... followed by the **month** part which can start with an optional 0 or 1
- 6 ... followed by a –
- 7 ... followed by the **year** part which must start with 19 or 20
- 8 ... followed by 2x digits
- 9 The character immediately after our date must be a non-digit

Possible solution

Below shows the 'evolution' of our pattern:

`\d\d-\d\d-\d\d\d\d`

`[0-3]` `\d-` `[01]` `\d-` `(19|20)` `\d\d`

`[0-3]` `\d-` `[01]` `\d-` `(19|20)` `\d{2}`

`[0-3]` `?` `\d-` `[01]` `?` `\d-` `(19|20)` `\d{2}`

`(?<=\D)` `[0-3]` `?` `\d-` `[01]` `?` `\d-` `(19|20)` `\d{2}` `(?=\D)`

`(?<=\D) [0-3]? \d- [01]? \d- (19|20) \d{2} (=?\D)`

- This pattern will successfully match against (for example):
 - TOP_SECRET_DATA_01-03-2020.CSV
 - TOP_SECRET_DATA_1-03-2020.CSV
 - TOP_SECRET_DATA_V2_20-3-2020.CSV
 - TOP_SECRET_DATA_5-4-1999_Adj.CSV

Possible solution

`(?<=\D) [0-3]? \d- [01]? \d- (19|20) \d{2} (=?\D)`

- Given the above pattern... **Which of the following would be successfully matched?**
 - **OUR_DATA_FILE_01-03-2120.CSV**
 - SKETCHY_INPUTS_V31-10-1999.CSV
 - **20-10-1999_NO_PEEKING.TXT**
 - **NO10_PARTY_INVITES_31_10_1999.CSV**

Using patterns within code

How do we actually use this pattern?

- Many programming languages provide *Regex* functionality.
 - In VBA via the Microsoft VBScript Regular Expressions 5.5 reference.
 - In Python via the `re` library.
 - In R via the `stringr` library.
- However, there *may* be some differences in the respective implementations.

How could we do this within R?

- Suppose we have some .CSV files within a folder.
- Let's assume that the variable `FOLDER_PATH` contains the path.
- Within our code, we can list all of the files in the folder...
- ...and then extract the dates from those files with a valid file name.

A worked example in R

```
tibble::tibble(  
  FILE_NAME =  
    fs::dir_ls(  
      path = FOLDER_PATH,  
      regexp = '(?i)CSV$' # ...another pattern!  
    ),  
  
  DATE_PART =  
    stringr::str_extract(  
      FILE_NAME,  
      pattern =  
      '(?<=\\D)[0-3]?\\d-[01]?\\d-(19|20)\\d{2}(?=\\D)'  
    )  
)
```

A worked example in R

Suppose our folder contained the following files:

20_10_1999_NO_PEEKING.CSV

N010_PARTY_INVITES_31-10-1999.CSV

OUR_DATA_FILE_01-03-2120.csv

SKETCHY_INPUTS_V31-10-1999.CSV

TOP_SECRET_DATA_01-03-2020.CSV

TOP_SECRET_DATA_1-03-2020.CSV

TOP_SECRET_DATA_5-4-1999_Adj.CSV

TOP_SECRET_DATA_V2_20-3-2020.CSV

NOT_A_CSV.TXT

A worked example in R

Below shows the output from our code on the above folder:

FILE_NAME	DATE_PART
20_10_1999_NO_PEEKING.CSV	NA
NO10_PARTY_INVITES_31-10-1999.CSV	NA
OUR_DATA_FILE_01-03-2120.csv	NA
SKETCHY_INPUTS_V31-10-1999.CSV	31-10-1999
TOP_SECRET_DATA_01-03-2020.CSV	01-03-2020
TOP_SECRET_DATA_1-03-2020.CSV	1-03-2020
TOP_SECRET_DATA_5-4-1999_Adj.CSV	5-4-1999
TOP_SECRET_DATA_V2_20-3-2020.CSV	20-3-2020

A worked example in R

- If we *pipe* our DATE_PART into lubridate::dmy, we can convert our extracted date into an actual date object that we can more easily work with.
- We can also use dplyr::filter to only retain those file names with valid dates.

```
tibble::tibble(  
  ...  
  DATE_PART =  
    ... %>%  
      lubridate::dmy()  
) %>%  
dplyr::filter(  
  !is.na(DATE_PART)  
)
```

A worked example in R

Our revised output is shown below:

FILE_NAME	DATE_PART
SKETCHY_INPUTS_V31-10-1999.CSV	1999-10-31
TOP_SECRET_DATA_01-03-2020.CSV	2020-03-01
TOP_SECRET_DATA_1-03-2020.CSV	2020-03-01
TOP_SECRET_DATA_5-4-1999_Adj.CSV	1999-04-05
TOP_SECRET_DATA_V2_20-3-2020.CSV	2020-03-20

Wrap up

What can I take away from this?

- That regex patterns can be used to match characters within some wider text.
- Using various elements within our pattern, we can refine what we are looking for.
- We are not just restricted to numbers!

What didn't we see?

- In our example above, we could have used *capture groups* to extract the individual day, month and year parts.
- ...these can also be used to make a pattern dependent on what has already been matched.
- Additional *quantifiers* such as `+` and `*`.
- Using *anchors* such as `^` and `$` to ensure that matches begin and/or finish at the start/end of a string respectively.
- Many other character classes such as `\w`, `\s` and `\n`.
- *Flags*; for example `(?i)` makes a pattern case-insensitive.

For those wanting to know more:

- <https://regexone.com/>
- <https://unicode-org.github.io/icu/userguide/strings/regexp.html>
- <https://medium.com/factory-mind/regex-tutorial-a-simple-cheatsheet-by-examples-649dc1c3f285>
- Speak to me.

That's it!

**Any questions?
... comments?**